
Review: Mitigating Memory Safety Vulnerabilities. Introduction to Cryptography

CMPSC 403 Fall 2021


October 5, 2021

Recall: Putting Together an Attack

1. Find a memory safety (e.g. buffer overflow) vulnerability
2. Write malicious shellcode at a known memory address
 - Mitigation: Address-space layout randomization
3. Overwrite the RIP with the address of the shellcode
 - Mitigation: Stack canaries
 - Mitigation: Pointer authentication
4. Return from the function
5. Begin executing malicious shellcode
 - Mitigation: Non-executable pages

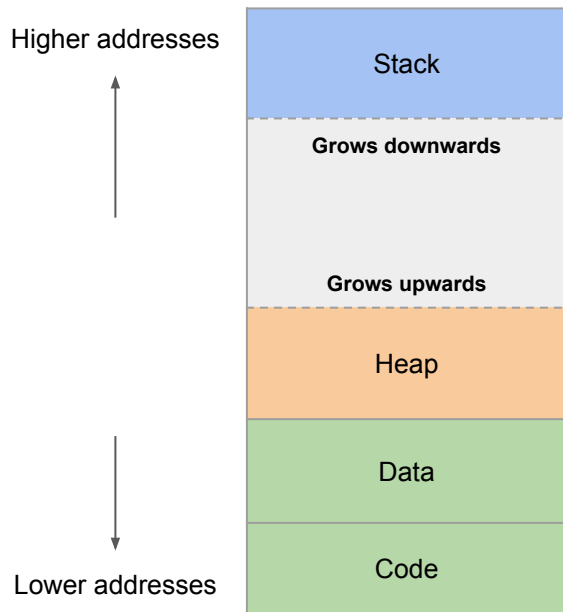
We can defend against memory safety vulnerabilities by making each of these steps more difficult (or impossible)!

Recall: x86 Memory Layout




In theory, x86 memory layout looks like this...

Recall: x86 Memory Layout




In theory, x86 memory layout looks like this...



...but in practice, it usually looks like this (mostly empty)!

Recall: x86 Memory Layout



Idea: Put each segment of memory in a different location each time the program is run

Address Space Layout Randomization

- **Address space layout randomization (ASLR):** Put each segment of memory in a different location each time the program is run
 - The attacker can't know where their shellcode will be because its address changes every time you run the program
- **ASLR can shuffle all four segments of memory**
 - Randomize the stack: Can't place shellcode on the stack without knowing the address of the stack
 - Randomize the heap: Can't place shellcode on the heap without knowing the address of the heap
 - Randomize the code

Combining Mitigations

Combining Mitigations

- Recall: We can use multiple mitigations together
 - Synergistic protection: one mitigation helps strengthen another mitigation
 - Force the attacker to find multiple vulnerabilities to exploit the program
 - Defense in depth
- Example: Combining ASLR and non-executable pages
 - An attacker can't write their own shellcode, because of non-executable pages
 - An attacker can't use existing code in memory, because they don't know the addresses of those code (ASLR)
- To defeat ASLR *and* non-executable pages, the attacker needs to find two vulnerabilities
 - First, find a way to leak memory and reveal the address randomization (defeat ASLR)
 - Second, find a way to write to memory and defeat non-executable pages

Enabling Mitigations

- Many mitigations are effectively free today (insignificant performance impact)
- The programmer sometimes has to manually enable mitigations
 - Example: Enable ASLR and non-executable pages when running a program
 - Example: Setting a flag to compile a program with stack canaries (secret value placed on the stack which changes every time the program is started to protect RIP)
- If the default is disabling the mitigation, the default will be chosen
 - Recall: Consider human factors!
 - Recall: Use fail-safe defaults!

Enabling Mitigations: Internet of Things



Qualys. Qualys Security Blog

[Link](#)

CVE-2021-3156: Heap-Based Buffer Overflow in Sudo (Baron Samedit)

Animesh Jain

January 26, 2021

The Qualys Research Team has discovered a heap overflow vulnerability in sudo, a near-ubiquitous utility available on major Unix-like operating systems. Any unprivileged user can gain root privileges on a vulnerable host using a default sudo configuration by exploiting this vulnerability.

Takeaway: Many (most?) IoT devices don't enable basic mitigations

Summary: Memory Safety Mitigations

- **Mitigation: Non-executable pages**
 - Make portions of memory either executable or writable, but not both
 - Defeats attacker writing shellcode to memory and executing it
 - Subversions
 - **Return-to-libc**: Execute an existing function in the C library
 - **Return-oriented programming (ROP)**: Create your own code by chaining together small gadgets in existing library code
- **Mitigation: Stack canaries**
 - Add a sacrificial value on the stack. If the canary has been changed, someone's probably attacking our system
 - Defeats attacker overwriting the RIP with address of shellcode
 - Subversions
 - An attacker can write around the canary
 - The canary can be leaked by another vulnerability (e.g. format string vulnerability)
 - The canary can be brute-forced by the attacker

Summary: Memory Safety Mitigations

- **Mitigation: Pointer authentication**
 - When storing a pointer in memory, replace the unused bits with a pointer authentication code (PAC). Before using the pointer in memory, check if the PAC is still valid
 - Defeats attacker overwriting the RIP (or any pointer) with address of shellcode
- **Mitigation: Address space layout randomization (ASLR)**
 - Put each segment of memory in a different location each time the program is run
 - Defeats attacker knowing the address of shellcode
 - Subversions
 - Leak addresses with another vulnerability
 - Brute-force attack to guess the addresses
- **Combining mitigations**
 - Using multiple mitigations usually forces the attacker to find multiple vulnerabilities to exploit the program (defense-in-depth)

Next: Cryptography

What is cryptography?

- Older definition: the study of secure communication over insecure channels
- Newer definition: provide rigorous guarantees on the security of data and computation in the presence of an attacker
- Modern cryptography involves a lot of math
 - We'll review any necessary material as they come up

Don't try this at home!

A Brief History of Cryptography

Cryptography by Hand: Caesar Cipher

- One of the earliest cryptographic schemes was the **Caesar cipher**
 - Used by Julius Caesar over 2,000 years ago
- Choose a key K randomly between 0 and 25
- To encrypt a plaintext message M :
 - Replace each letter in M with the letter K positions later in the alphabet
 - If $K = 3$, plaintext DOG becomes GRJ
- To decrypt an encrypted ciphertext C :
 - Replace each letter in C with the letter K positions earlier in the alphabet
 - If $K = 3$, ciphertext GRJ becomes DOG

$K = 3$			
M	C	M	C
A	D	N	Q
B	E	O	R
C	F	P	S
D	G	Q	T
E	H	R	U
F	I	S	V
G	J	T	W
H	K	U	X
I	L	V	Y
J	M	W	Z
K	N	X	A
L	O	Y	B
M	P	Z	C



Cryptography by Hand: Substitution Cipher

- A better cipher: create a mapping of each character to another character.
 - Example: A = N, B = Q, C = L, D = Z, etc.
 - Unlike the Caesar cipher, the shift is no longer constant!
- Key generation algorithm: $\text{KeyGen}()$
 - Generate a random, one-to-one mapping of characters
- Encryption algorithm: $\text{Enc}(K, M)$
 - Map each letter in M to the output according to the mapping K
- Decryption algorithm: $\text{Dec}(K, C)$:
 - Map each letter in C to the output according to the **reverse** of the mapping K

K			
M	C	M	C
A	N	N	G
B	Q	O	P
C	L	P	T
D	Z	Q	A
E	K	R	J
F	R	S	O
G	V	T	D
H	U	U	I
I	E	V	C
J	S	W	F
K	B	X	M
L	W	Y	X
M	Y	Z	H


Cryptography by Machines: Enigma

- A mechanical encryption machine used by the Germans in WWII



Cryptography by Machines: Enigma

- KeyGen():
 - Choose rotors, rotor orders, rotor positions, and plugboard settings
 - 158,962,555,217,826,360,000 possible keys
- Enc(K, M) and Dec(K, C):
 - Input the rotor settings K into the Enigma machine
 - Press each letter in the input, and the lampboard will light up the corresponding output letter
 - Encryption and decryption are the same algorithm!
- Germans believed that Enigma was an “unbreakable code”



Cryptography by Computers

- The modern era of cryptography started after WWII, with the work of Claude Shannon
- “New Directions in Cryptography” (1976) showed how number theory can be used in cryptography
 - Its authors, Whitfield Diffie and Martin Hellman, won the Turing Award in 2015 for this paper

Takeaways

- Cryptography started with paper-and-pencil algorithms (Caesar cipher)
- Then cryptography moved to machines (Enigma)
- Finally, cryptography moved to computers (which we're about to study)

Definitions

Meet Alice, Bob, Janyl, and Jada

- Alice and Bob: The main characters trying to send messages to each other over an insecure communication channel
- Janyl: An **eavesdropper** who can read any data sent over the channel
- Jada: A **manipulator** who can read and modify any data sent over the channel

Meet Alice, Bob, Janyl, and Jada

- We often describe cryptographic problems using a common cast of characters
- One scenario:
 - Alice wants to send a message to Bob.
 - However, Janyl is going to *eavesdrop* on the communication channel.
 - How does Alice send the message to Bob without Janyl learning about the message?
- Another scenario:
 - Bob wants to send a message to Alice.
 - However, Jada is going to *tamper* with the communication channel.
 - How does Bob send the message to Alice without Jada changing the message?

Three Goals of Cryptography

- In cryptography, there are three main properties desired in data
- **Confidentiality:** An adversary cannot **read** our messages.
- **Integrity:** An adversary cannot **change** our messages without being detected.
- **Authenticity:** I can prove that this message came from the person who claims to have written it.
 - Integrity and authenticity are closely related properties...
 - Before I can prove that a message came from a certain person, I have to prove that the message wasn't changed!
 - ... but they're not identical properties

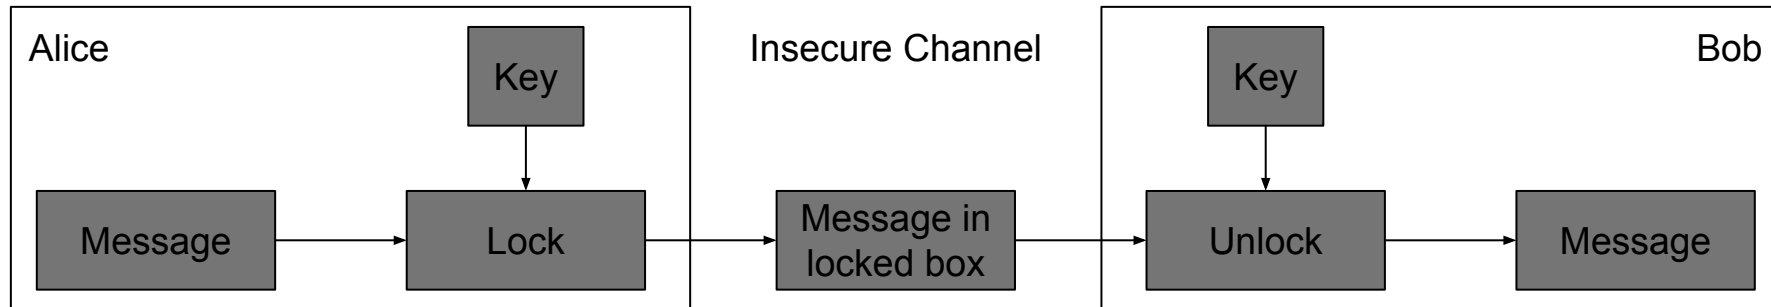
—

Keys

- The most basic building block of any cryptographic scheme: The **key**
- We can use the key in our algorithms to secure messages
- Two models of keys:
 - **Symmetric key model:** Alice and Bob both know the value of the same secret key.
 - **Asymmetric key model:** Everybody has two keys, a secret key and a public key.
 - Example: RSA encryption

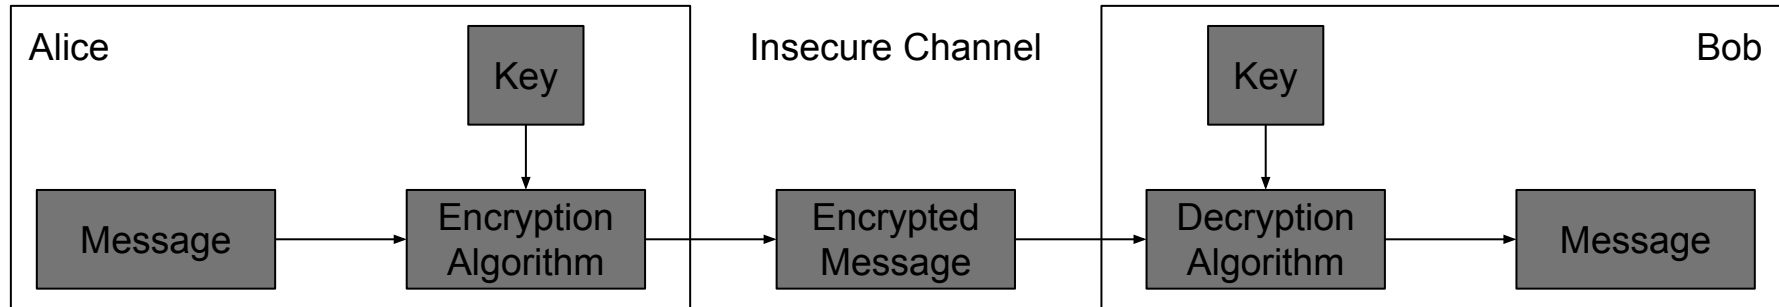
Confidentiality

- **Confidentiality:** An adversary cannot **read** our messages.
- **Analogy:** Locking and unlocking the message
 - Alice uses the key to lock the message in a box
 - Alice sends the message (locked in the box) over the insecure channel
 - Janyl sees the locked box, but cannot access the message without the key
 - Bob receives the message (locked in the box) and uses the key to unlock the message



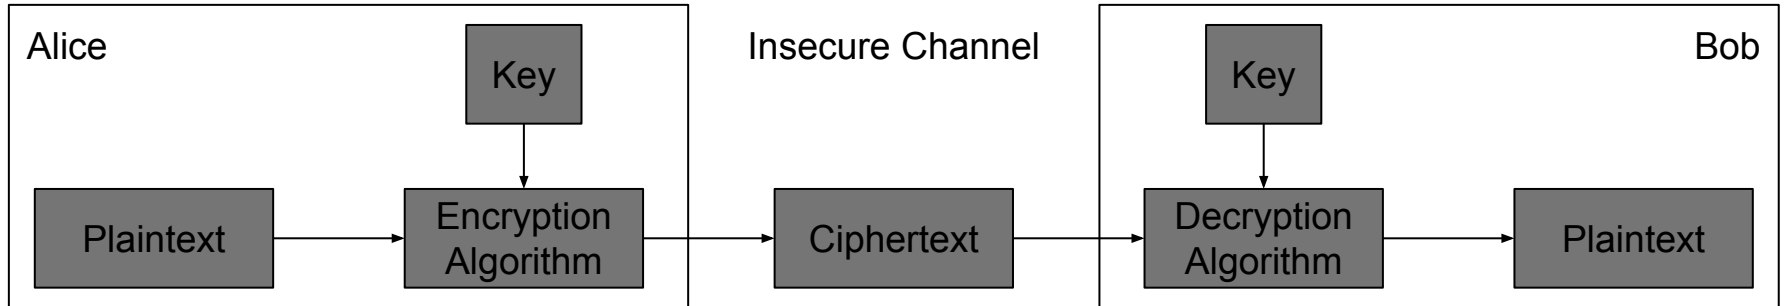
Confidentiality

- **Confidentiality:** An adversary cannot **read** our messages.
- Schemes provide confidentiality by **encrypting** messages
 - Alice uses the key to **encrypt** the message: change the message into a scrambled form
 - Alice sends the encrypted message over the insecure channel
 - Janyl sees the encrypted message, but cannot figure out the original message without the key
 - Bob receives the encrypted message and uses the key to **decrypt** the message back into its original form



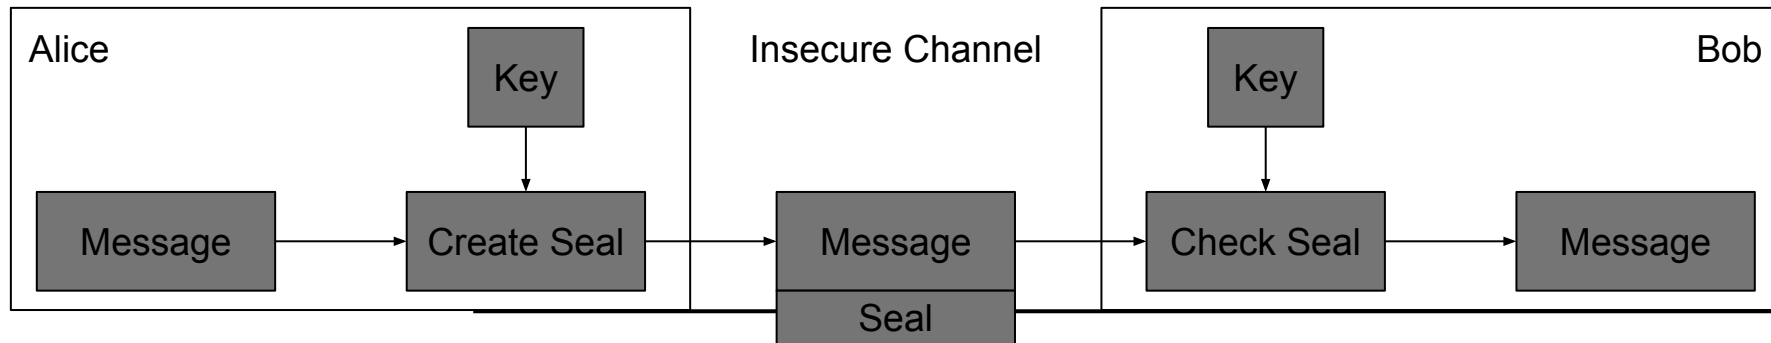
Confidentiality

- **Plaintext:** The original message
- **Ciphertext:** The encrypted message



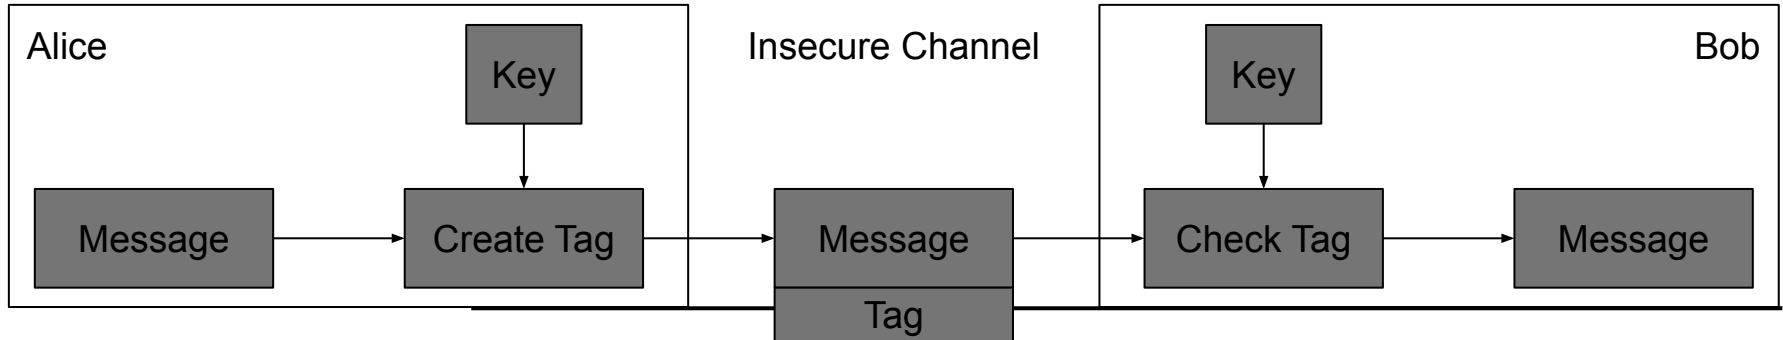
Integrity (and Authenticity)

- **Integrity:** An adversary cannot **change** our messages without being detected.
- **Analogy:** Adding a seal on the message
 - Alice uses the key to add a special seal on the message (e.g. puts tape on the envelope)
 - Alice sends the message and the seal over the insecure channel
 - If Jada tampers with the message, she'll break the seal (e.g. break the tape on the envelope)
 - Without the key, Jada cannot create her own seal
 - Bob receives the message and the seal and checks that the seal has not been broken



Integrity (and Authenticity)

- **Integrity:** An adversary cannot **change** our messages without being detected.
- Schemes provide integrity by adding a **tag** or **signature** on messages
 - Alice uses the key to generate a special tag for the message
 - Alice sends the message and the tag over the insecure channel
 - If Jada tampers with the message, the tag will no longer be valid
 - Bob receives the message and the tag and checks that the tag is still valid
- More on integrity in a future lecture



Threat Models

- What if Janyl can do more than eavesdrop?
- Real-world schemes are often vulnerable to more sophisticated attackers, so cryptographers have created more sophisticated threat models too
- Some threat models for analyzing confidentiality:

	Can Janyl trick Alice into encrypting messages of Janyl's choosing?	Can Janyl trick Bob into decrypting messages of Janyl's choosing?
Ciphertext-only		
Chosen-plaintext	✓	
Chosen-ciphertext		✓
Chosen plaintext-ciphertext	✓	✓

Cryptographic Hash Function: Definition

- Hash function: $H(M)$
 - Input: *Arbitrary* length message M
 - Output: *Fixed* length, n -bit hash
 - Sometimes written as $\{0, 1\}^* \rightarrow \{0, 1\}^n$
- Properties
 - **Correctness:** Deterministic
 - Hashing the same input always produces the same output
 - **Efficiency:** Efficient to compute
 - **Security:** One-way-ness (“preimage resistance”)
 - **Security:** Collision-resistance
 - **Security:** random/unpredictability, no predictable patterns for how changing the input affects the output
 - Changing 1 bit in the input causes the output to be completely different
 - Also called “random oracle” assumption

Hash Function: Intuition


- A hash function provides a fixed-length “fingerprint” over a sequence of bits
- Example: Document comparison
 - If Alice and Bob both have a 1 GB document, they can both compute a hash over the document and (securely) communicate the hashes to each other
 - If the hashes are the same, the files must be the same, since they have the same “fingerprint”
 - If the hashes are different, the files must be different

Hash Function: One-way-ness

- **Informal:** Given an output y , it is infeasible to find **any** input x such that $H(x) = y$
- **Intuition:** Here's an output. Can you find an input that hashes to this output?
 - Note: The adversary just needs to find *any* input, not necessarily the input that was actually used to generate the hash

Is this function one-way?

- The constant function $H(x) = 1$
 - No, because an attacker can output any x , and that leads to 1. It does not have the original x some challenger thought about
- Take fruit and make a smoothie: yes



Ref:
https://computersciencewiki.org/index.php/One-way_function

- Block cipher, for k random and secret: E
 - Yes, if an attacker can invert E_k , then an attacker can distinguish it from a random permutation which breaks the security of the block cipher